

File Format Benchmark - Avro, JSON, ORC, & Parquet

Owen O'Malley

owen@hortonworks.com

[@owen_omalley](https://twitter.com/owen_omalley)

September 2016

Who Am I?

- ◆ Worked on Hadoop since Jan 2006
- ◆ MapReduce, Security, Hive, and ORC
- ◆ Worked on different file formats
 - Sequence File, RCFile, ORC File, T-File, and Avro requirements

Goal

- ◆ Seeking to discover unknowns
 - How do the different formats perform?
 - What could they do better?
 - Best part of open source is looking inside!
- ◆ Use real & diverse data sets
 - Over-reliance on similar datasets leads to weakness
- ◆ Open & reviewed benchmarks

The File Formats

Avro

- ◆ Cross-language file format for Hadoop
- ◆ Schema evolution was primary goal
- ◆ Schema segregated from data
 - Unlike Protobuf and Thrift
- ◆ Row major format

JSON

- ◆ Serialization format for HTTP & Javascript
- ◆ Text-format with MANY parsers
- ◆ Schema completely integrated with data
- ◆ Row major format
- ◆ Compression applied on top

ORC

- ◆ Originally part of Hive to replace RCFile
 - Now top-level project
- ◆ Schema segregated into footer
- ◆ Column major format with stripes
- ◆ Rich type model, stored top-down
- ◆ Integrated compression, indexes, & stats

Parquet

- ◆ Design based on Google's Dremel paper
- ◆ Schema segregated into footer
- ◆ Column major format with stripes
- ◆ Simpler type-model with logical types
- ◆ All data pushed to leaves of the tree
- ◆ Integrated compression and indexes

Data Sets

NYC Taxi Data

- ◆ Every taxi cab ride in NYC from 2009
 - Publically available
 - <http://tinyurl.com/nyc-taxi-analysis>
- ◆ 18 columns with no null values
 - Doubles, integers, decimals, & strings
- ◆ 2 months of data – 22.7 million rows



Github Logs

- ◆ All actions on Github public repositories
 - Publically available
 - <https://www.githubarchive.org/>
- ◆ 704 columns with a lot of structure & nulls
 - Pretty much the kitchen sink
- ◆ 1/2 month of data – 10.5 million rows

Finding the Github Schema

- ◆ The data is all in JSON.
- ◆ No schema for the data is published.
- ◆ We wrote a JSON schema discoverer.
 - Scans the document and figures out the type
- ◆ Available at <https://github.com/hortonworks/hive-json>
- ◆ Schema is huge (12k)

Sales

◆ Generated data

- Real schema from a production Hive deployment
- Random data based on the data statistics

◆ 55 columns with lots of nulls

- A little structure
- Timestamps, strings, longs, booleans, list, & struct

◆ 25 million rows

Storage costs

Compression

◆ Data size matters!


- Hadoop stores all your data, but requires hardware
- Is one factor in read speed

◆ ORC and Parquet use RLE & Dictionaries

◆ All the formats have general compression

- ZLIB (GZip) – tight compression, slower
- Snappy – some compression, faster


Taxi Size



Taxi Size Analysis

- ◆ Don't use JSON
- ◆ Use either Snappy or Zlib compression
- ◆ Avro's small compression window hurts
- ◆ Parquet Zlib is smaller than ORC
 - Group the column sizes by type


Parquet & ORC Taxi Column Sizes





Taxi Size Analysis

- ◆ ORC did better than expected
 - String columns have small cardinality
 - Lots of timestamp columns
 - No doubles 😊
- ◆ Need to revalidate results with original
 - Improve random data generator
 - Add non-smooth distributions



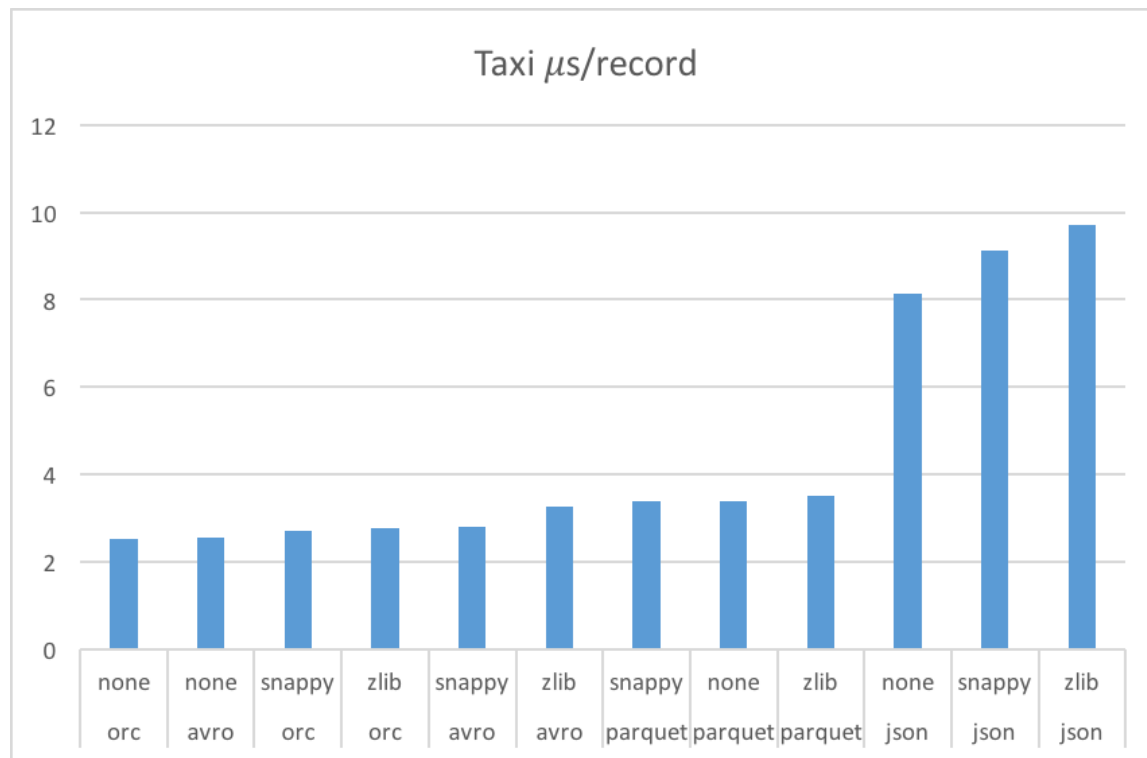
Github Size Analysis

- ◆ Surprising win for JSON and Avro
 - Worst when uncompressed
 - Best with zlib
- ◆ Many partially shared strings
 - ORC and Parquet don't compress across columns
- ◆ Need to investigate shared dictionaries.

Use Cases


Full Table Scans

- ◆ Read all columns & rows
- ◆ All formats except JSON are splittable
 - Different workers do different parts of file



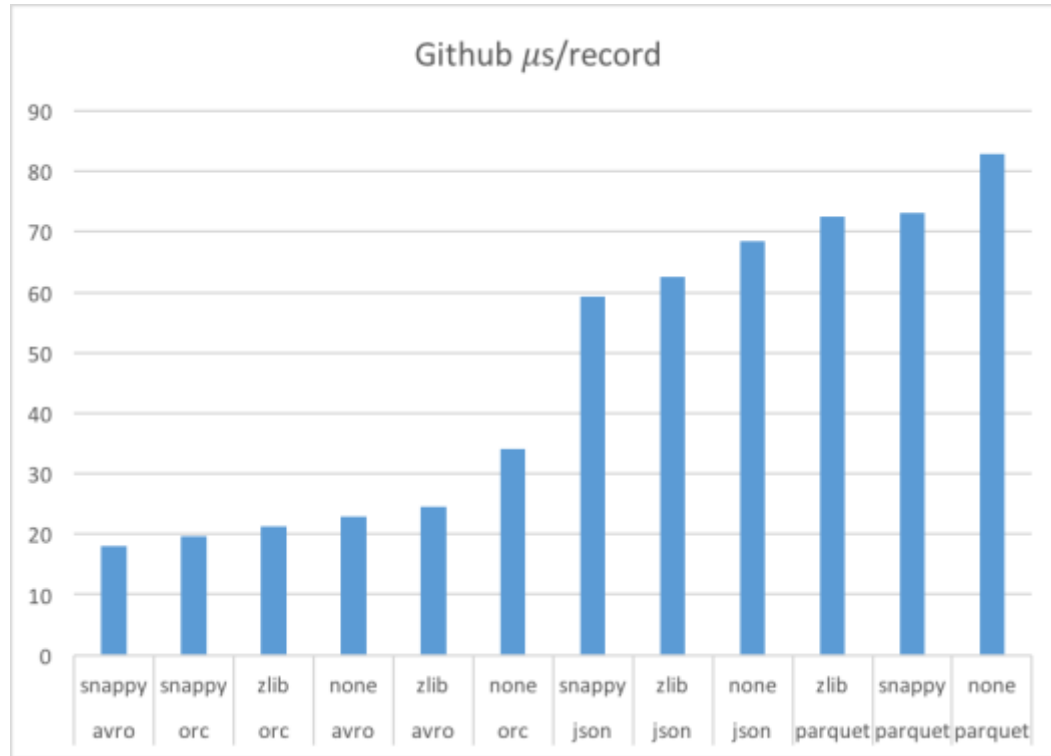
Taxi Read Performance Analysis

- ◆ JSON is very slow to read
 - Large storage size for this data set
 - Needs to do a **LOT** of string parsing
- ◆ Tradeoff between space & time
 - Less compression is sometimes faster



Sales Read Performance Analysis

- ◆ Read performance is dominated by format
 - Compression matters less for this data set
 - Straight ordering: ORC, Avro, Parquet, & JSON
- ◆ Garbage collection is important
 - ORC 0.3 to 1.4% of time
 - Avro < 0.1% of time
 - Parquet 4 to 8% of time



Github Read Performance Analysis

- ◆ Garbage collection is critical
 - ORC 2.1 to 3.4% of time
 - Avro 0.1% of time
 - Parquet 11.4 to 12.8% of time
- ◆ A lot of columns needs more space
 - Suspect that we need bigger stripes
 - Rows/stripe - ORC: 18.6k, Parquet: 88.1k

Column Projection

- ◆ Often just need a few columns
 - Only ORC & Parquet are columnar
 - Only read, decompress, & deserialize some columns

Dataset	format	compression	us/row	projection	Percent time
github	orc	zlib	21.319	0.185	0.87%
github	parquet	zlib	72.494	0.585	0.81%
sales	orc	zlib	1.866	0.056	3.00%
sales	parquet	zlib	12.893	0.329	2.55%
taxi	orc	zlib	2.766	0.063	2.28%
taxi	parquet	zlib	3.496	0.718	20.54%

Projection & Predicate Pushdown

- ◆ Sometimes have a filter predicate on table
 - Select a superset of rows that match
 - Selective filters have a huge impact
- ◆ Improves data layout options
 - Better than partition pruning with sorting
- ◆ ORC has added optional bloom filters

Metadata Access

◆ ORC & Parquet store metadata

- Stored in file footer
- File schema
- Number of records
- Min, max, count of each column

◆ Provides $O(1)$ Access

Conclusions

Recommendations

- ◆ Disclaimer – **Everything** changes!
 - Both these benchmarks and the formats will change.
- ◆ Don't use JSON for processing.
- ◆ If your use case needs column projection or predicate push down:
 - ORC or Parquet

Recommendations

- ◆ For complex tables with common strings
 - Avro with Snappy is a good fit (w/o projection)
- ◆ For other tables
 - ORC with Zlib or Snappy is a good fit
- ◆ Tweet benchmark suggestions to [@owen_omalley](#)

Fun Stuff

- ◆ Built open benchmark suite for files
- ◆ Built pieces of a tool to convert files
 - Avro, CSV, JSON, ORC, & Parquet
- ◆ Built a random parameterized generator
 - Easy to model arbitrary tables
 - Can write to Avro, ORC, or Parquet

Remaining work

- ◆ Extend benchmark with LZO and LZ4.
- ◆ Finish predicate pushdown benchmark.
- ◆ Add C++ reader for ORC, Parquet, & Avro.
- ◆ Add Presto ORC reader.

Thank you!

Twitter: @owen_omalley

Email: owen@hortonworks.com